


**YΣ19 Artificial Intelligence II (Deep Learning for  
Natural Language Processing)  
Fall Semester 2021  
Homework 2  
20% of the course mark  
Announced: November 25, 2021 Due: December 16,  
2021 at 23:59**

1. Compute the gradient of the cross entropy loss function with respect to its logits. See case (7), page 4 of the [matrix calculus](#) notes of the Stanford class CS222n. Remember that we have defined cross entropy loss in our lecture on softmax regression. Note that the notes from Stanford use an alternative notation which you are expected to follow.

(1/10 marks)

2. Apply backpropagation to the following computational graph to compute all the gradients with respect to the inputs  $x, m, b$  and  $y^*$ . Assume the loss function is  $MSE(y, y^*) = (y - y^*)^2$  where  $y^*$  is the true label and  $y$  is the predicted output from the neural network. Note also that  $ReLU(x) = \max(0, x)$ .



(1/10 marks)

3. Develop a vaccine sentiment classifier with 3 classes (neutral, anti-vax and pro-vax) using feed-forward neural networks and the datasets of the previous homework provided on e-class. We expect you to experiment with at least two models and compare them using F1 score, recall and precision. For the development of the models, you are expected to experiment at least with the following hyperparameters:

- the number of hidden layers and the number of their units
- the activation functions
- the loss function
- the optimizer

Furthermore, in at least one of your models, you must use pre-trained word embedding vectors using [GloVe](#).

For your best model, plot the loss vs epochs and the [ROC curve](#) and explain what you see. Also, compare your best model with the model you obtained using softmax regression in Homework 1. Your solution should be implemented in PyTorch and we expect your report to be well-documented.

(8/10 marks)